AI’s cyber risk arrives as a thousand cuts

送信者: EG Geo-technology 受信日時: 2026-11-12 05:57

【メール原文】

A single cyber catastrophe is unlikely, but many smaller attacks can amount to systemic damage. October brought a wave of high-profile cyber incidents‌‌‌‌‌‌‌‌‌‌‌‌‌‌‌‌

Geo-technology
Customize Email Preferences

The cyber risk of AI arrives as a thousand cuts

11 Nov 2026 15:56 EST Discover more on the Portal

Listen now · 5 min

*

Criminals and hacktivists gain the most from AI’s offensive cyber capabilities, while states remain constrained by geopolitics, making a single catastrophic attack unlikely.

*

Security spending is rising but concentrating in well-funded firms, leaving resource-constrained organizations exposed.

*

Many smaller attacks, or one that spreads through shared infrastructure, could have systemic consequences, but the diffuse threat will slow any policy response.

A single cyber catastrophe is unlikely, but many smaller attacks can amount to systemic damage. October brought a wave of high-profile cyber incidents, from breaches at South Korean banks to the $6.3 billion Binance theft, the largest crypto theft on record. Media coverage has focused on AI’s role in these attacks and the prospect of a single AI-enabled strike that could take down a power grid or a financial exchange. For most organizations, the more relevant risk is their own exposure. AI’s main effect is speed and scale: It can find existing vulnerabilities faster than most defenders can fix them. Vulnerabilities remain easiest to find in small, poorly defended, and less-regulated systems, favoring ordinary cybercriminals and hacktivists.

AI is unlikely to change the calculus of state-backed groups. State-sponsored and state-affiliated groups remain the world’s best-resourced offensive cyber organizations. AI adds capability, but for these groups, the binding constraints tend to be political: their stake in stability and the geopolitical cost of exposure. Those constraints hold in peacetime, but escalating conflict would weaken them. Rogue states with little reputation to protect, such as North Korea, behave more like cybercriminals. They keep attacking, but primarily for financial gain. The Binance theft, attributed to a North Korean group, was an elaborate financial scheme, not an attempt to trigger broader disruption.

Open-weight models put frontier attack capability in criminal hands within months of each release. They trail the closed frontier by four to eight months. Because anyone can download, modify, and run them privately, developers cannot monitor or limit misuse. For actors unconcerned about reputational risk, AI relieves talent and resource constraints. In the South Korean bank attacks, agentic AI tools let a single

financially motivated operator breach at least five banks in a short span. Attackers in the Cairn campaign used AI to breach more than 1,200 websites and over 100 companies. Whether they seek profit or pursue a cause, these actors have little incentive to bring down entire systems.

Most defenders cannot close vulnerabilities as fast as AI can find them. Cybersecurity has always been asymmetric: Defenders must secure every entry point, while attackers need only one. AI widens the gap by letting attackers probe more systems in less time. Even among teams deploying AI defensively, few have the resources and operational agility to patch faster than attackers can discover vulnerabilities. Decades-old weaknesses in legacy systems are exposed again, adding to the backlog of already stretched security teams.

Everyone will face attacks, but only well-resourced organizations will be able to absorb them. Global security spending will grow 12% to more than $240 billion in 2026, but the median US security budget is not growing. Well-funded organizations will continue to invest in cutting-edge tooling, hire world-class talent, and benefit from partnerships with frontier AI labs. Because they detect intrusions early and recover through redundant systems, these organizations can treat attacks as a cost of doing business. Smaller and public-sector organizations have no such cushion. For those running legacy systems with small teams and limited budgets, a single breach can halt operations.

Shared infrastructure can transmit damage from the least-defended organizations into the broader system. No single criminal or hacktivist needs to aim for a meltdown for one to happen. When many attackers exploit the same weakness at once, or when a weakness in a decentralized system is exploited, damage can spread across organizations and sectors without anyone coordinating it. In October, AI agents from startup Cathedral, testing under Anthropic’s bug bounty program, failed for days to breach Anthropic directly, then reached its sales, human resources, and IT systems in under 12 hours through third-party business software. If a frontier AI lab can be reached through its suppliers, most organizations likely can be too. Shared infrastructure such as the Domain Name System (DNS) can turn local failures into systemic ones, allowing a single attack to spread far beyond its intended target. Without a defining incident to force action, regulation will lag the threat.

Noah Fehr
Analyst, Geo-technology
fehr@eurasiagroup.net
Heather West
Practice Head, Geo-technology
west@eurasiagroup.net
Gerard DiPippo
Director, Global Macro
dipippo@eurasiagroup.net
EG Geo-technology
geo-technology@eurasiagroup.net

Now on iOS and Android
Get the Eurasia Group app
Read anywhere · Receive notifications · Reach out to analysts
App Store

Google Play

© 2026 Eurasia Group. All rights reserved. This material was produced by Eurasia Group for use by the recipient. This is intended as general background research and is not intended to constitute advice on any particular commercial investment or trade matter or issue and should not be relied upon for such purposes. It is not to be made available to any person other than the recipient. No part of this publication may be reproduced, stored in a retrieval system, or transmitted in any form or by any means, electronic or otherwise, without the prior consent of Eurasia Group.
X-HG-DocID cmwcl8kg4xd1z07n2s5j06939

Opt Out | Customize Email Preferences